

# Memory Architecture I

Michael C. Hackett  
Assistant Professor, Computer Science

Community  
College  
*of* Philadelphia

# Lecture Topics


- Sequential Circuits
  - Clocks
  - SR Latch
  - Flip-Flops
    - SR Flip-Flop
    - D Flip-Flop
    - JK Flip Flop
  - Registers
- Memory Hierarchy
- Memory Technologies
  - RAM
  - ROM
- Cache Memory
  - Direct Mapping
  - N-way Associative Set Mapping
  - Fully Associative Set Mapping

# Sequential Circuits

- The combinational logic circuits we've seen have the following limitations:
  - They have no memory capability
  - The output changes as soon as the input changes
- **Sequential logic circuits** maintain a state- The circuit's output is determined by both:
  - The input it currently has
  - The input it has received over time


# Clocks

- In most cases, the components in a sequential circuit need to be synchronized
  - The components in the circuit must all update their states at the same time
- This is achieved with a **clock signal**
  - A 1-bit signal that alternates between 1 and 0 at regular intervals



# Clocks

- A **clock cycle** (*“tick”*) is when the clock transitions from 0 to 1 and back to 0
- A **clock period** is the *time* it takes to complete 1 clock cycle



# Clocks

- The **clock frequency** (also called the **clock rate**) is the number of clock periods in some amount of (*wall clock*) time
  - **Wall clock time** is the seconds, minutes, hours, etc. that we experience
  - Time ticks by on a wall clock by the second, whereas these clocks may tick thousands (or millions or billions) of times per one second
- Clock frequency is measured in  $\frac{\text{cycle}}{\text{second}}$  or Hertz (abbreviated Hz)

# Clocks

- Clock frequency is the inverse of the clock period (and vice versa)

$$Period = \frac{seconds}{cycle} = \frac{1}{\frac{cycles}{second}} = \frac{1}{Frequency} \quad \text{Seconds per cycle}$$

$$Frequency = \frac{cycles}{second} = \frac{1}{\frac{seconds}{cycle}} = \frac{1}{Period} \quad \text{Cycles per second}$$

# Clocks

- Example: What is the clock rate of a processor with a clock period of 250 ps (picoseconds)?
  - $250ps = 250 * 10^{-12}seconds = 0.00000000025s$
- $Period = \frac{250 \times 10^{-12}seconds}{1\ cycle} = \frac{0.00000000025\ seconds}{1\ cycle} = 0.00000000025 \frac{seconds}{cycle}$
- $Frequency = \frac{1}{Period} = \frac{1}{\frac{250 \times 10^{-12}seconds}{1\ cycle}} = \frac{250 \times 10^{-12}cycles}{1\ second} = \frac{4,000,000,000\ cycles}{1\ second}$   
 $= 4,000,000,000 \frac{cycles}{second} = 4.0\ GHz$



# Clocks



- Example: What is the clock period of a processor with a clock rate of 3.8GHz?

- $Frequency = 3,800,000,000 \frac{cycles}{second} = \frac{3,800,000,000 cycles}{1 second}$

- $Period = \frac{1}{Frequency} = \frac{1}{\frac{3,800,000,000 cycles}{1 second}} = \frac{1 second}{3,800,000,000 cycles}$   
 $= 2.63 \times 10^{-10} \frac{seconds}{cycle} = 263 \times 10^{-12} \frac{seconds}{cycle}$   
 $= 263 \frac{picoseconds}{cycle}$


# Clocks

- Abstraction of a clock:



# SR Latch


- A **latch** is a digital component that stores 1 bit of information
- An SR Latch has two inputs (Set and Reset) and two outputs ( $Q$  and its complement,  $\bar{Q}$ )



# SR Latch

- **Feedback**

- The output of the top NOR is one input to the bottom NOR
- The output of the bottom NOR is one input to the top NOR



# SR Latch


- Recall that, for a NOR gate, the output is 0 if either input is a 1

$x$	$y$	$x \text{ NOR } y$
0	0	1
<b>0</b>	<b>1</b>	<b>0</b>
<b>1</b>	<b>0</b>	<b>0</b>
<b>1</b>	<b>1</b>	<b>0</b>

# SR Latch

- We'll start with the second row of the latch's truth table:
  - This is the Reset state

$S$	$R$	$\bar{Q}$	$Q$	State
0	0			
0	1	1	0	Reset ( $Q = 0$ )
1	0			
1	1			




- R is 1, meaning the bottom NOR must have an output (Q) of 0
- S is 0 and the output of the bottom NOR was 0, so the top NOR must have an output ( $\bar{Q}$ ) of 1
  - This is a valid state since Q and  $\bar{Q}$  are complements

# SR Latch

- Next is the third row of the latch's truth table:
  - This is the Set state

$S$	$R$	$\bar{Q}$	$Q$	State
0	0			
0	1	1	0	Reset ( $Q = 0$ )
1	0	0	1	Set ( $Q = 1$ )
1	1			




- S is 1, meaning the top NOR must have an output ( $\bar{Q}$ ) of 0
- R is 0, and the output of the top NOR was 0, so the bottom NOR must have an output (Q) of 1
  - This is a valid state since Q and  $\bar{Q}$  are complements

# SR Latch

- Next is the fourth row of the latch's truth table:
  - This is the Unknown state

$S$	$R$	$\bar{Q}$	$Q$	State
0	0			
0	1	1	0	Reset ( $Q = 0$ )
1	0	0	1	Set ( $Q = 1$ )
1	1	0	0	Unknown



- S is 1 and R is 1 meaning both NOR gates must have an output of 0
- This doesn't make sense since  $Q$  and  $\bar{Q}$  are supposed to be complements



# SR Latch

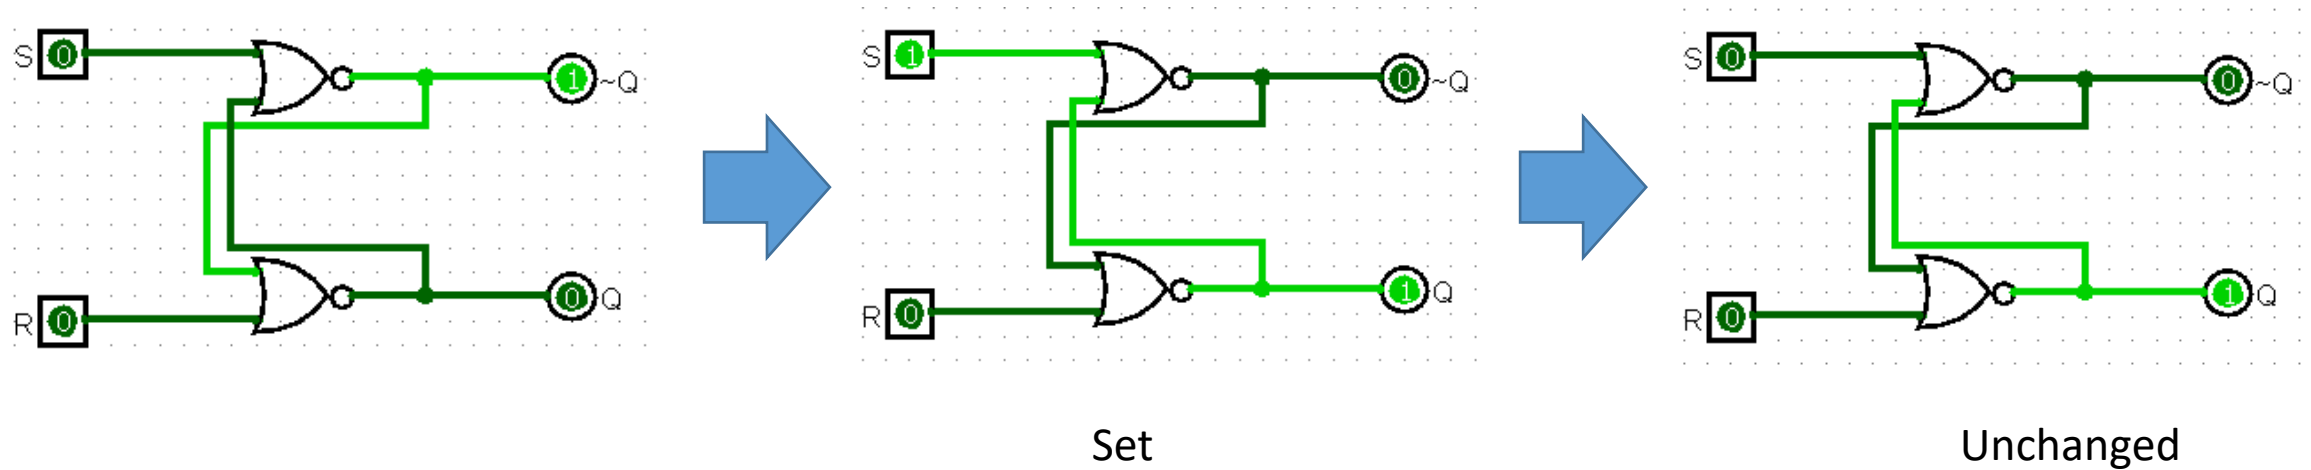
- Back to the first row of the latch's truth table:
  - This is the Unchanged state

$S$	$R$	$\bar{Q}$	$Q$	State
0	0	$\bar{Q}$	$Q$	Unchanged
0	1	1	0	Reset ( $Q = 0$ )
1	0	0	1	Set ( $Q = 1$ )
1	1	0	0	Unknown

- $S$  is 0 and  $R$  is 0, the current state will depend on the previous state
  - If the previous state was the set state, the current state will still be in the set state
  - If the previous state was the reset state, the current state will still be in the reset state



# SR Latch

- Turning Q from 0 (Reset) to 1 (Set)
  - $S = 0, S = 1, S = 0$




# SR Latch

- Turning Q from 1 (Set) to 0 (Reset)
  - R = 0, R = 1, R = 0




Reset



Unchanged

# SR Flip-Flop


- An extension of the SR Latch is the SR Flip-Flop.
- The S and R inputs are each and'ed with a clock signal.
  - The state can only be changed when the clock signal is 1



$S$	$R$	Clock	$\bar{Q}$	$Q$	State
0	0	1	$\bar{Q}$	$Q$	Unchanged
0	1	1	1	0	Reset ( $Q = 0$ )
1	0	1	0	1	Set ( $Q = 1$ )
1	1	1	0	0	Unknown
X	X	0	$\bar{Q}$	$Q$	Unchanged

# D Flip-Flop


- A D Flip-Flop has one input (D) and two outputs ( $Q$  and its complement,  $\bar{Q}$ )
- The input is and'ed with a clock signal prior to the NOR gates.
  - The state can only be changed when the clock signal is 1



$D$	Clock	$\bar{Q}$	$Q$	State
0	1	1	0	Reset ( $Q = 0$ )
1	1	0	1	Set ( $Q = 1$ )
$X$	0	$\bar{Q}$	$Q$	Unchanged

# JK Flip-Flop


- In a JK Flip-Flop, the input is and'ed with a clock signal *and* an output prior to the NOR gates.
  - The state can only be changed when the clock signal is 1




$J$	$K$	Clock	$\bar{Q}$	$Q$	State
0	0	1	$\bar{Q}$	$Q$	Unchanged
0	1	1	1	0	Reset ( $Q = 0$ )
1	0	1	0	1	Set ( $Q = 1$ )
1	1	1	$Q$	$\bar{Q}$	Toggle
X	X	0	$\bar{Q}$	$Q$	Unchanged

# Flip-Flops


- Abstractions of Flip-Flops:



SR Flip Flop



D Flip Flop



JK Flip Flop

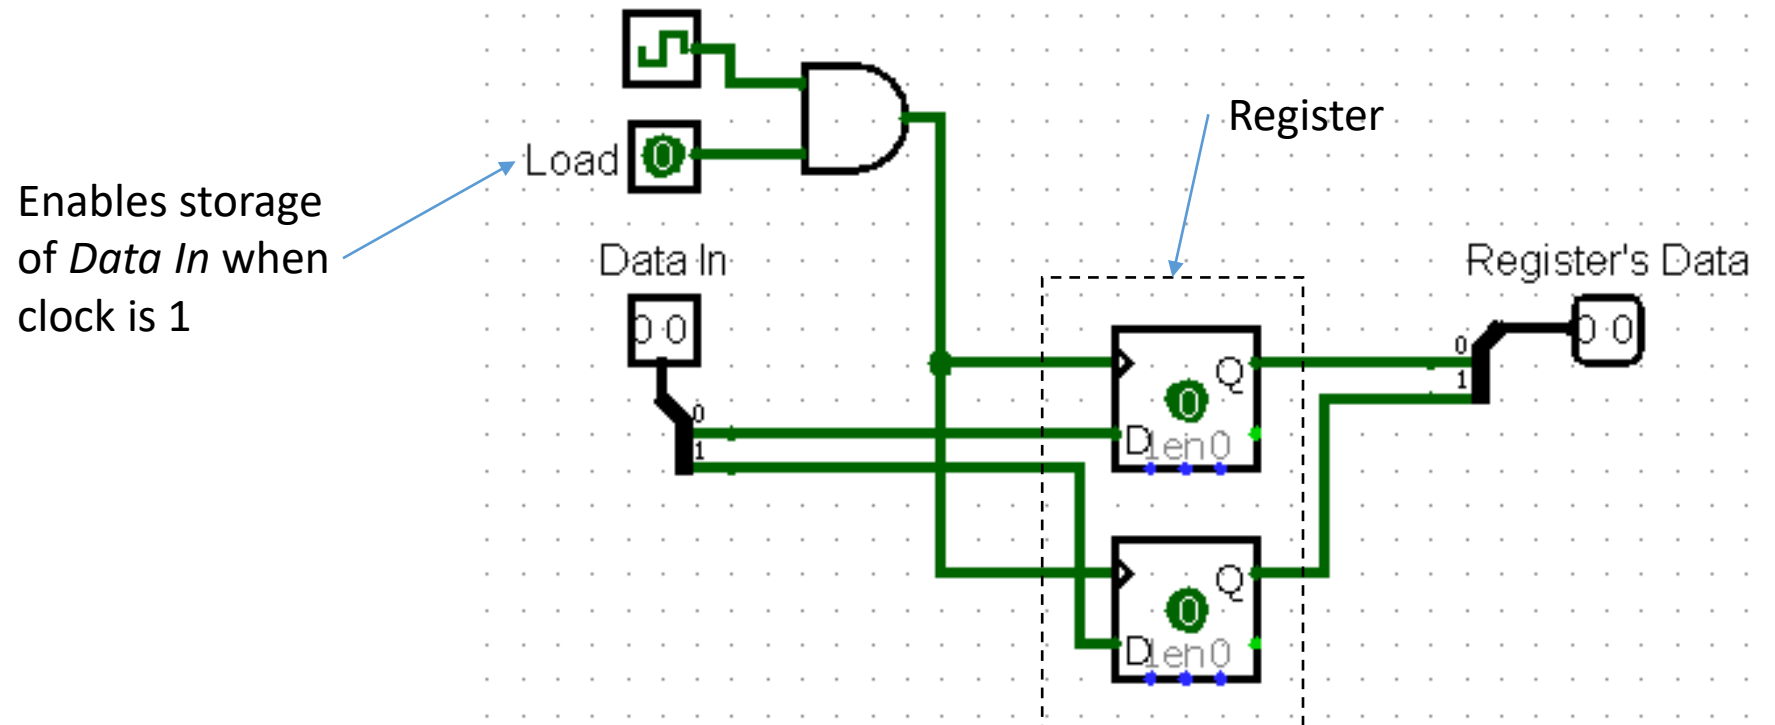
# Registers

- By now, we are familiar with the use of registers from assembly programming to temporarily store data.
- Since flip-flops store 1 bit of information, we can create a register from a series of flip-flops
  - 4 flip-flops for a 4-bit register, 32 flip-flops for a 32-bit register, etc.



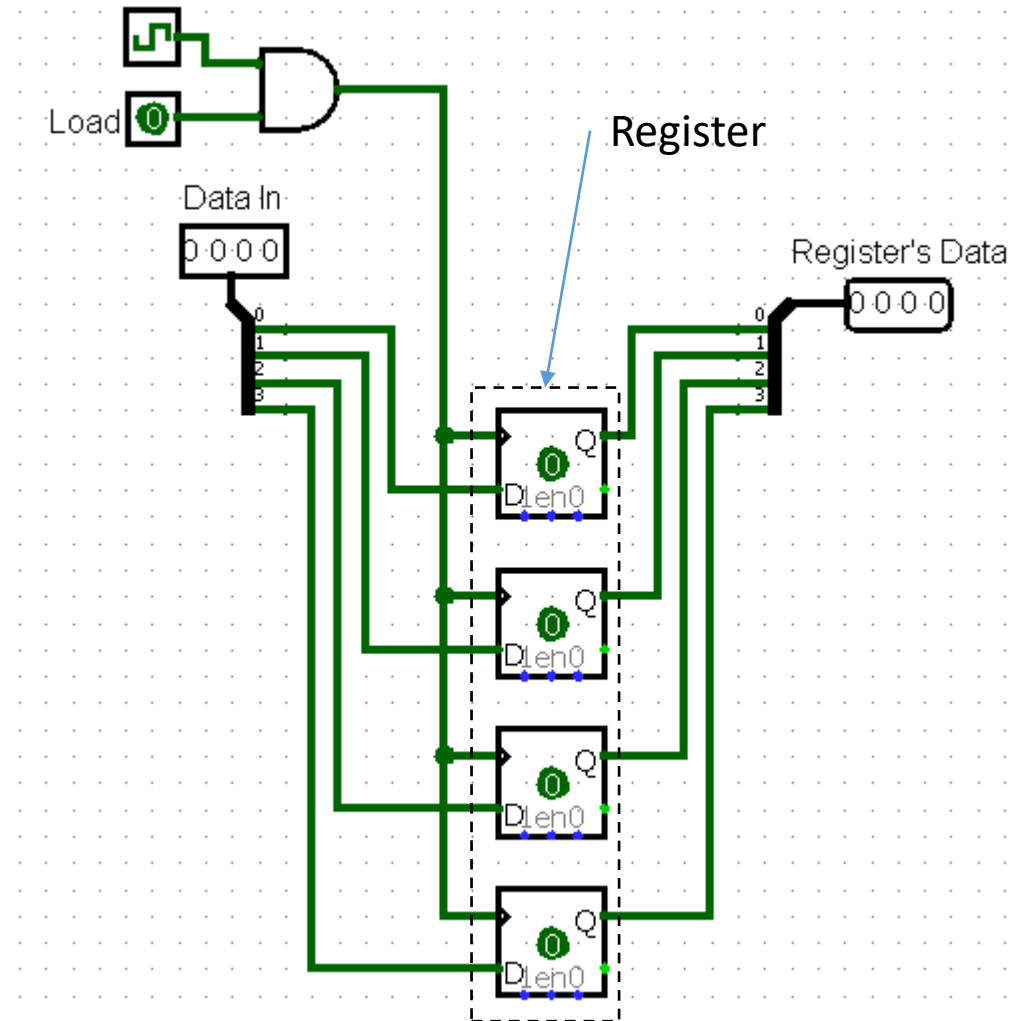
# Registers

- A 2-bit register:



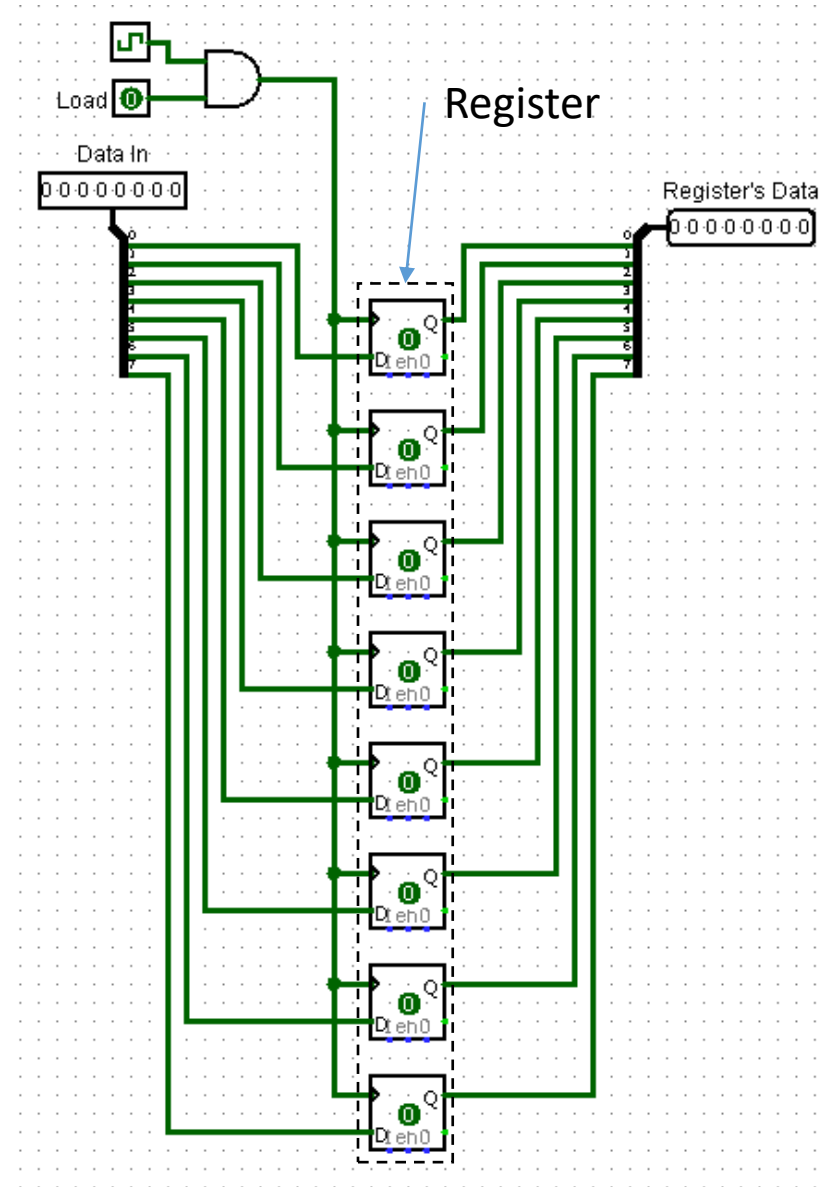
# Registers

- A 4-bit register:



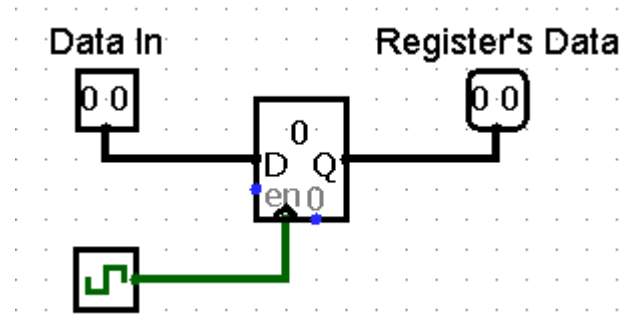
# Registers

- An 8-bit register:

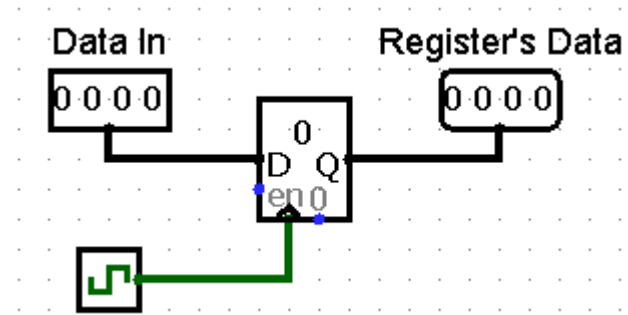


# Registers

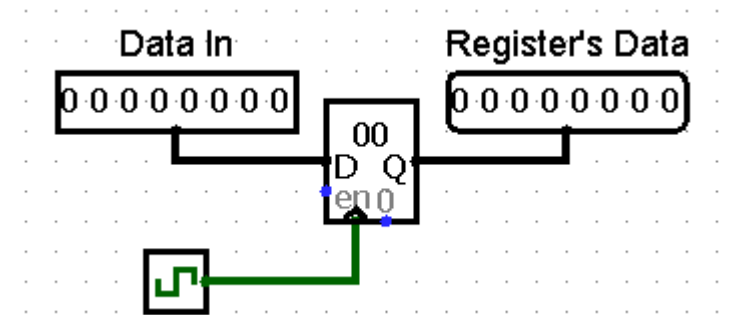
- Abstractions of Registers:



2-Bit Register



4-Bit Register



8-Bit Register

# Memory Hierarchy

- The **Memory Hierarchy** is the structured levels of a computer system's memory.
  - As the distance from the processor increases, the size and the access time increases

- |                      |                    |
|----------------------|--------------------|
| 1. Registers         | (Fastest/Smallest) |
| 2. Cache Memory      |                    |
| 3. Main Memory       |                    |
| 4. Secondary Storage | (Slowest/Largest)  |

# Memory Hierarchy

- In the first tier are *registers*, which is memory used for storing data currently in use by the CPU.
- Registers have the fastest access time, but they are limited in number.

# Memory Hierarchy

- In the second tier is *cache memory*, which is memory used for storing data the CPU has recently used.
- The cache memory is built into the CPU and has a much larger capacity than the limited number of registers.
  - Cache memory is still limited in space, as only so much can fit in the CPU.

# Memory Hierarchy

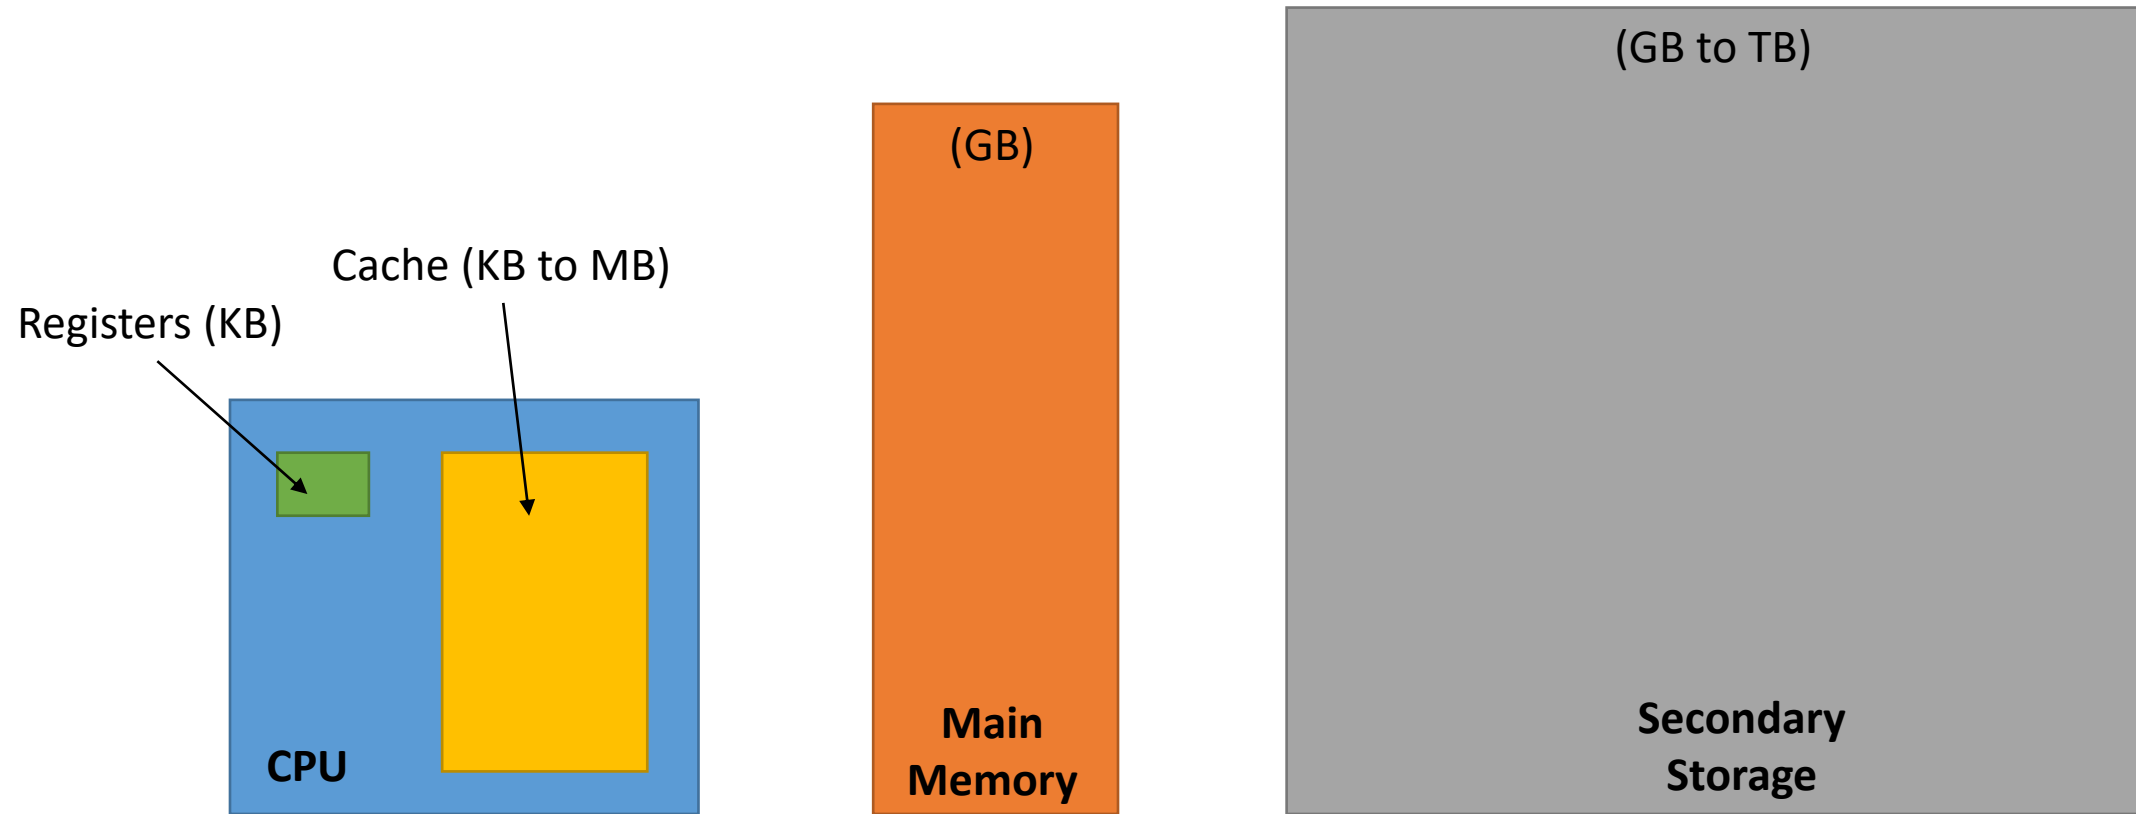
- In the third tier is *main memory*, which is memory used for storing data the CPU doesn't immediately need.
- Main Memory is often referred to as the system's *RAM*
  - In this course, RAM will indicate a specific type of memory technology.
- Much larger capacity than the CPU's cache memory.



# Memory Hierarchy

- In the fourth tier is *secondary storage*, which is memory used for storing data long term.
  - Magnetic Disks and Tape, Flash Memory, Optical Disks, etc.
- It is the memory with the slowest access time but has the greatest capacities.
- Secondary storage use **non-volatile** memory technologies.
  - The data stored in these technologies remains stored even when the system's power is turned off.
- Registers, cache and main memory use **volatile** memory technologies.
  - The data stored in these technologies are lost when the system's power is turned off.

# Memory Hierarchy



# Memory Technologies

- There are a variety of memory technologies used for secondary storage.
  - We'll discuss them in the next lecture
- This lecture focuses on the more fundamental types of memory technologies: RAM and ROM

# Memory Technologies


- In older computer systems, reels of magnetic tape were used for main memory and secondary storage
- To access data, the reels were spun forward and reverse
  - *Sequential Access*
- And yes, tape is still a thing for secondary storage
  - Cost efficient means of backing up data



# RAM

- **RAM** (**R**andom **A**ccess **M**emory) is a type of volatile memory that does not retain its stored bits when the system is powered off.
- Data stored in RAM can be accessed at random using an address
- Consists of  $N$  words of  $M$  bits each ( $N \times M$  memory)
  - Each word has a unique address
- For example, a  $4096 \times 8$  memory:
  - 4096 8-bit words (32768 bits total)
  - 4096 unique addresses (0 through 4095)


# RAM



# RAM

- **Static Random Access Memory (SRAM)** and **Dynamic Random Access Memory (DRAM)** are types of RAM built from an array of memory cells
- Each cell can store one bit of information
  - SRAM memory cell: Uses transistors and a loop of not gates
  - DRAM memory cell: Uses transistors and a capacitor
- Each cell is connected to a word enable line and data line
  - The **word enable line** allows the cell's value to be changed
  - The **data line** is the data to be stored in the cell


# RAM – SRAM Cell






# RAM

Memory Cells







# RAM

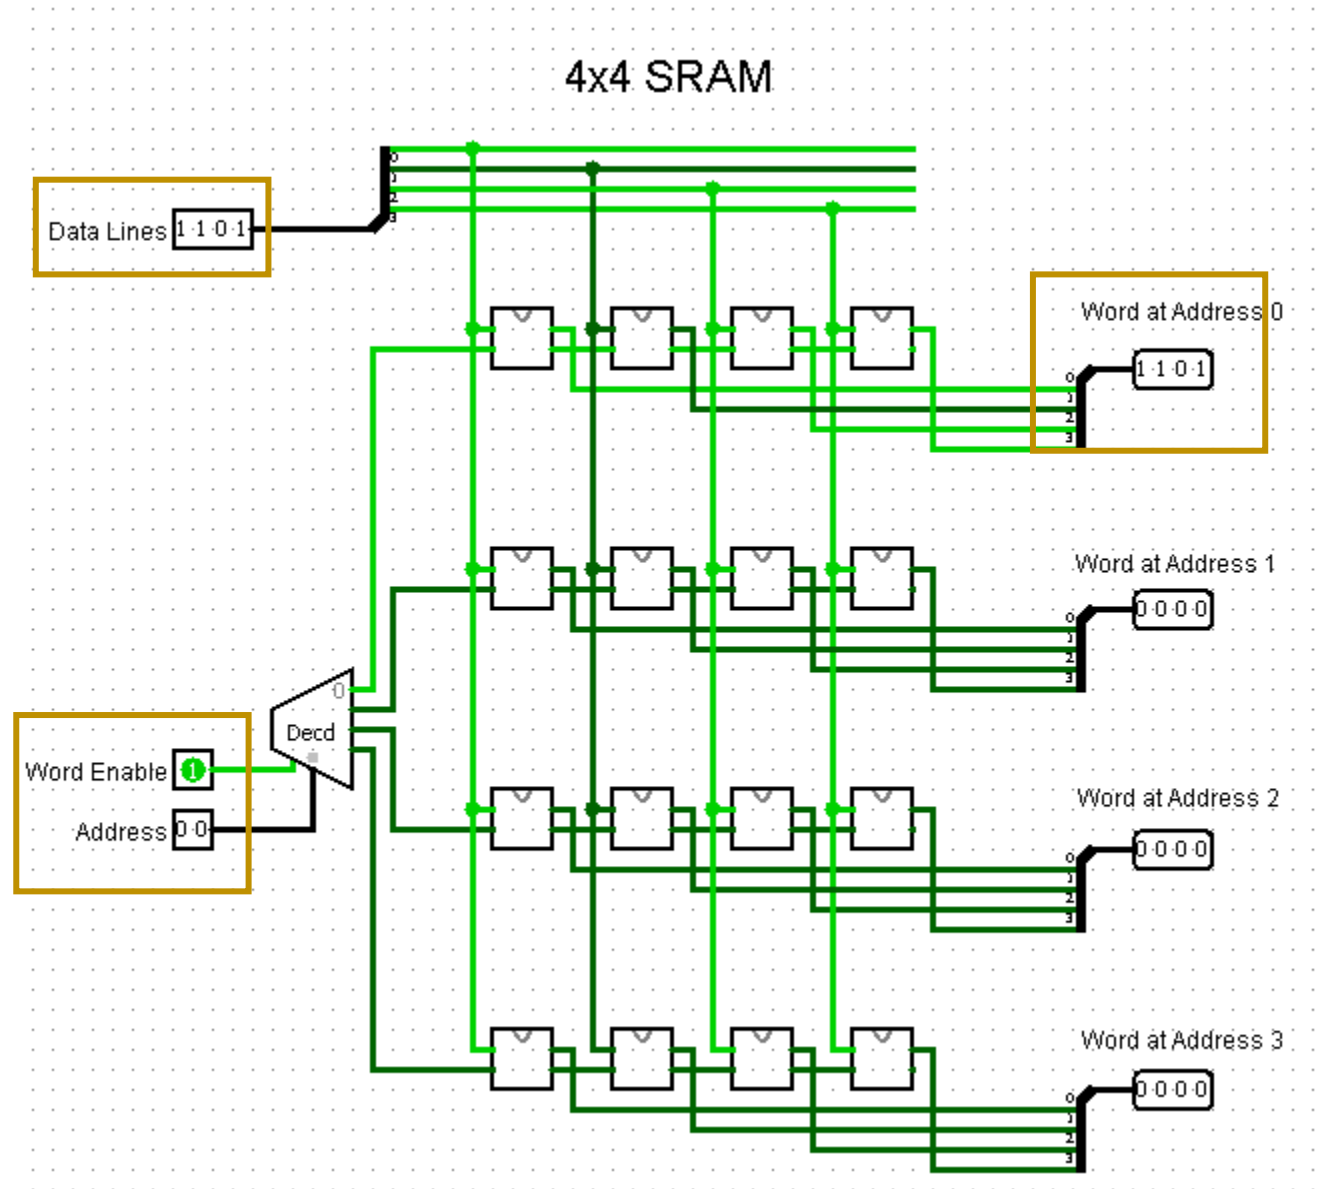
- To select a word, the word's address is decoded to enable that row of cells
  - The “Word Enable” here is controlling if the decoder is enabled
    - Each output is a word enable to a row of cells
  - The “Address” is the input to the decoder



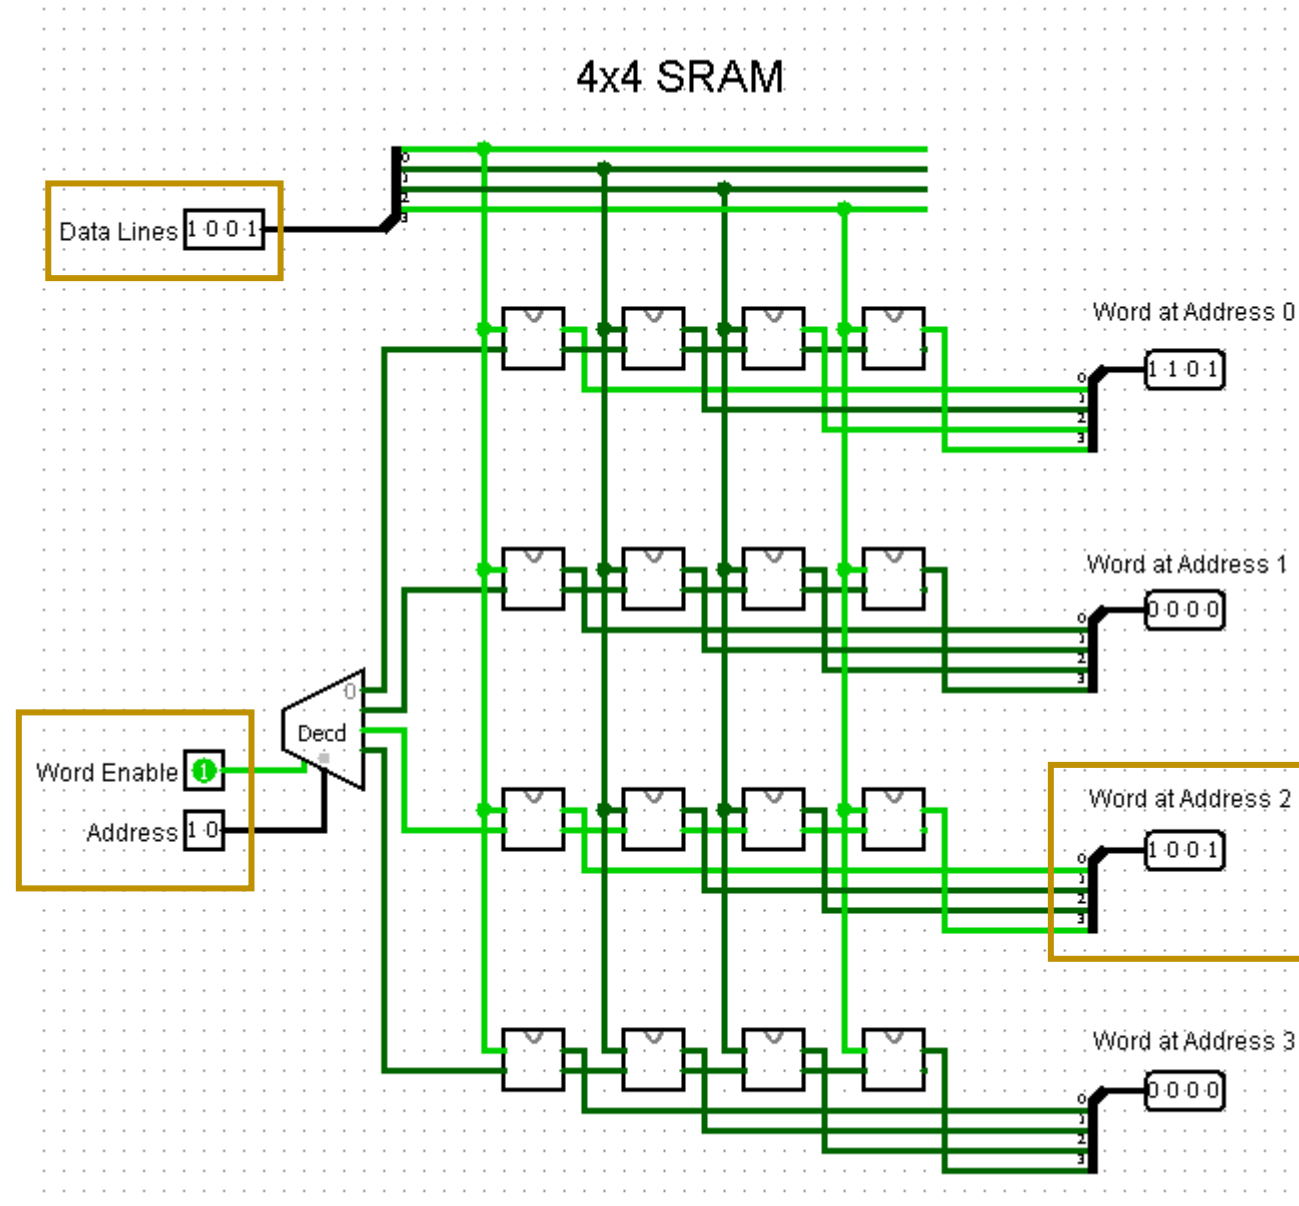
# RAM



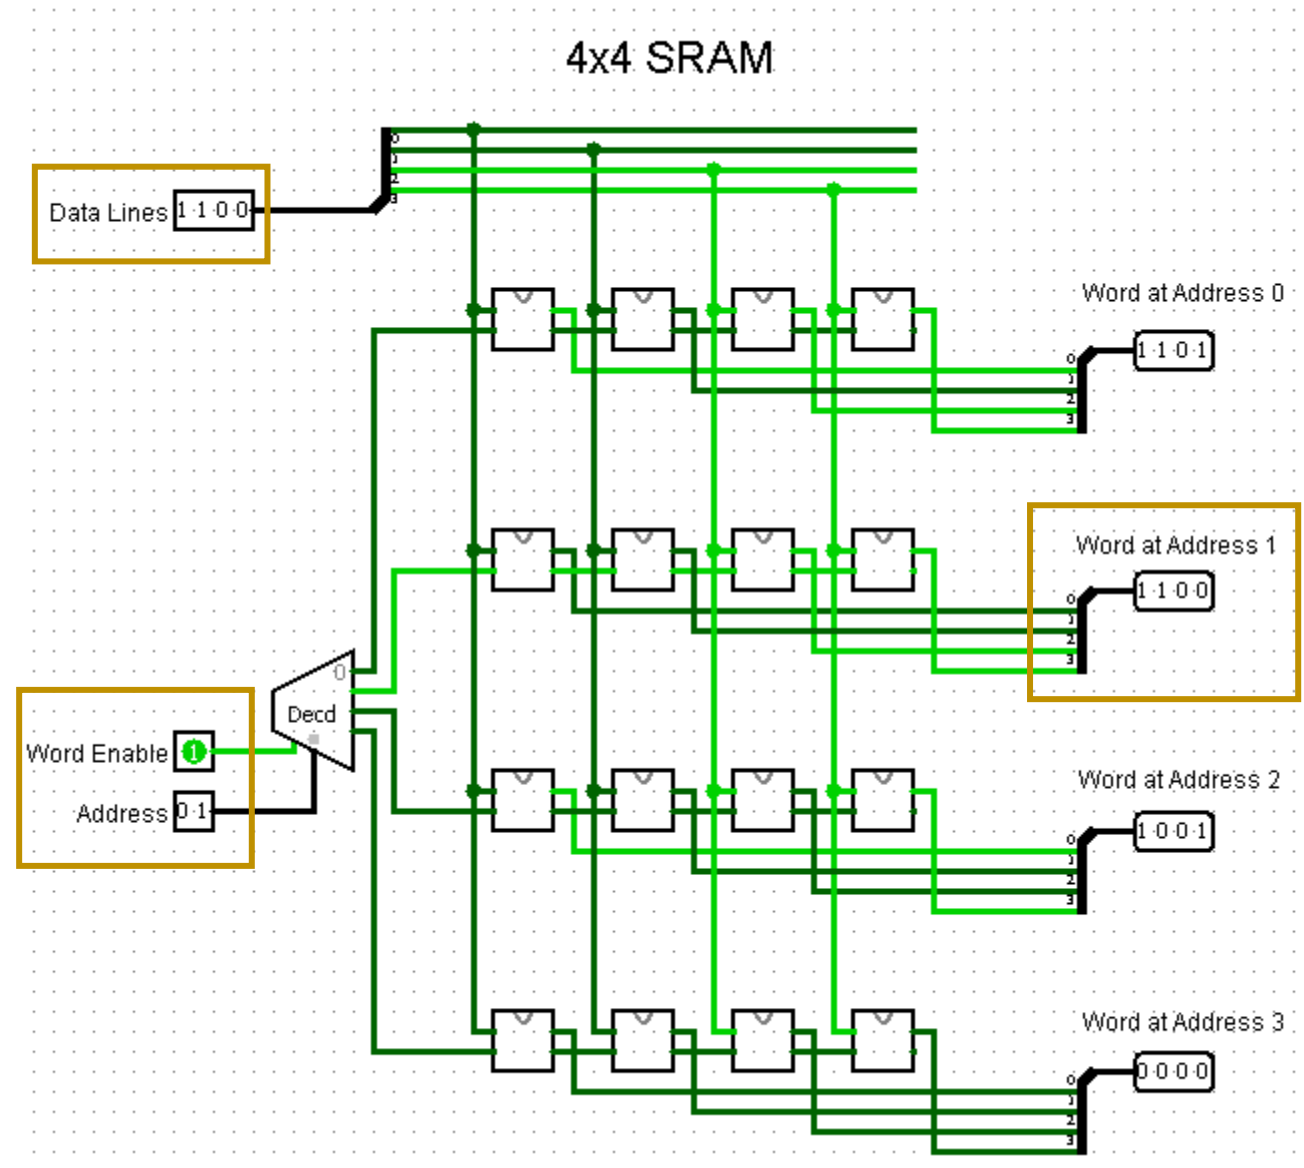
# RAM



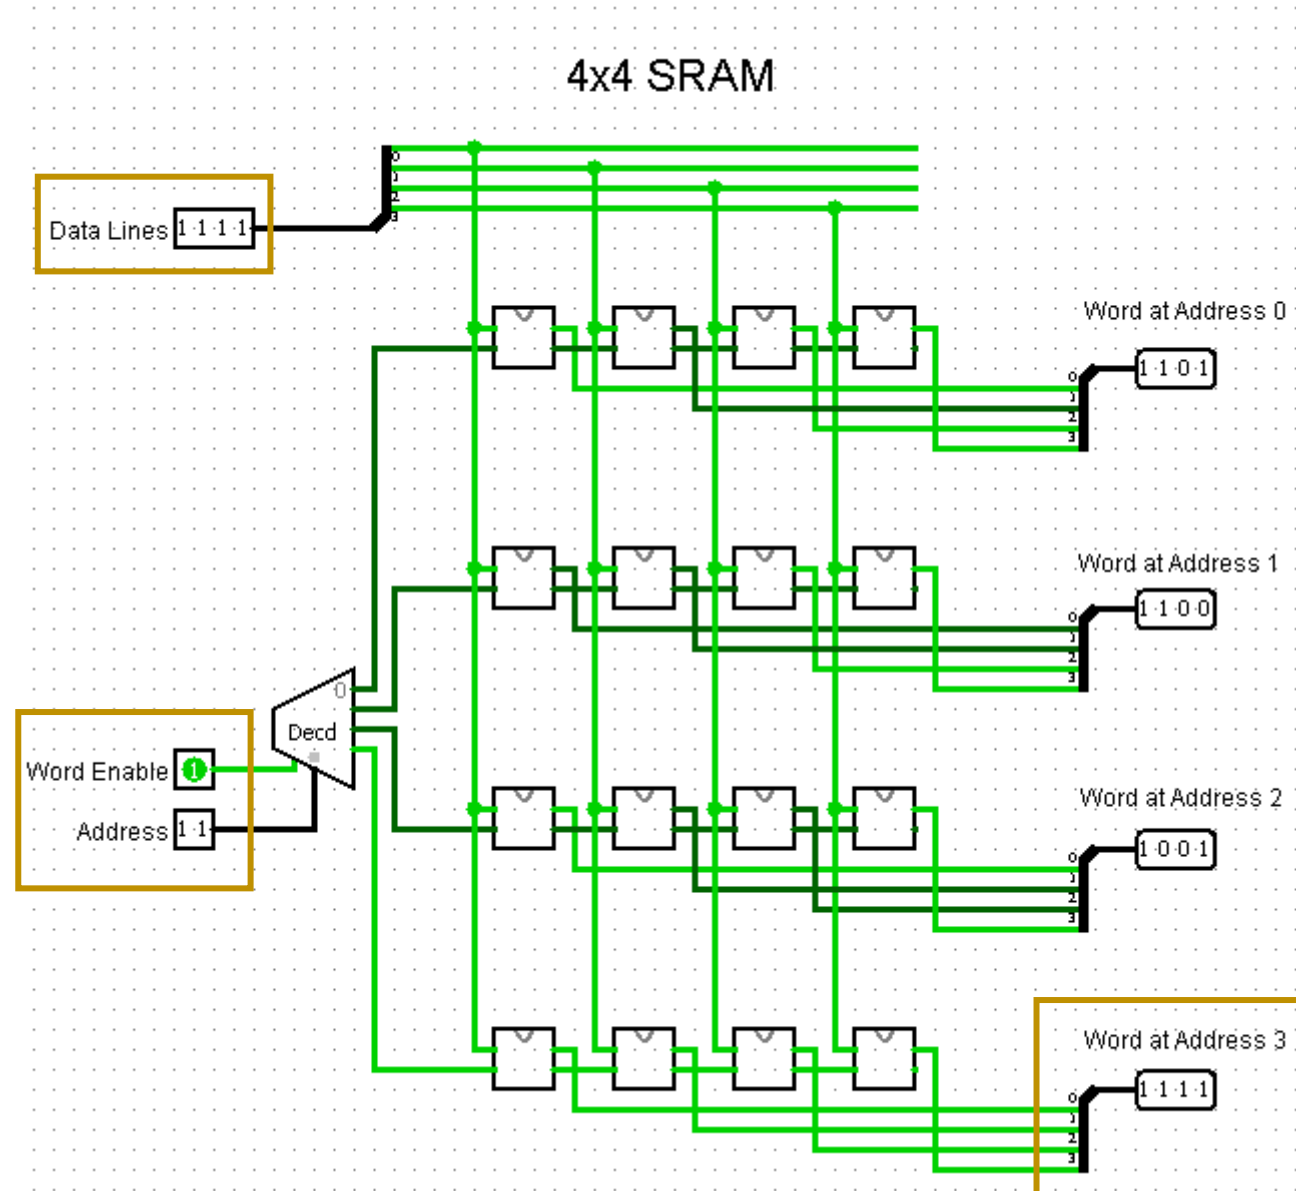
# RAM



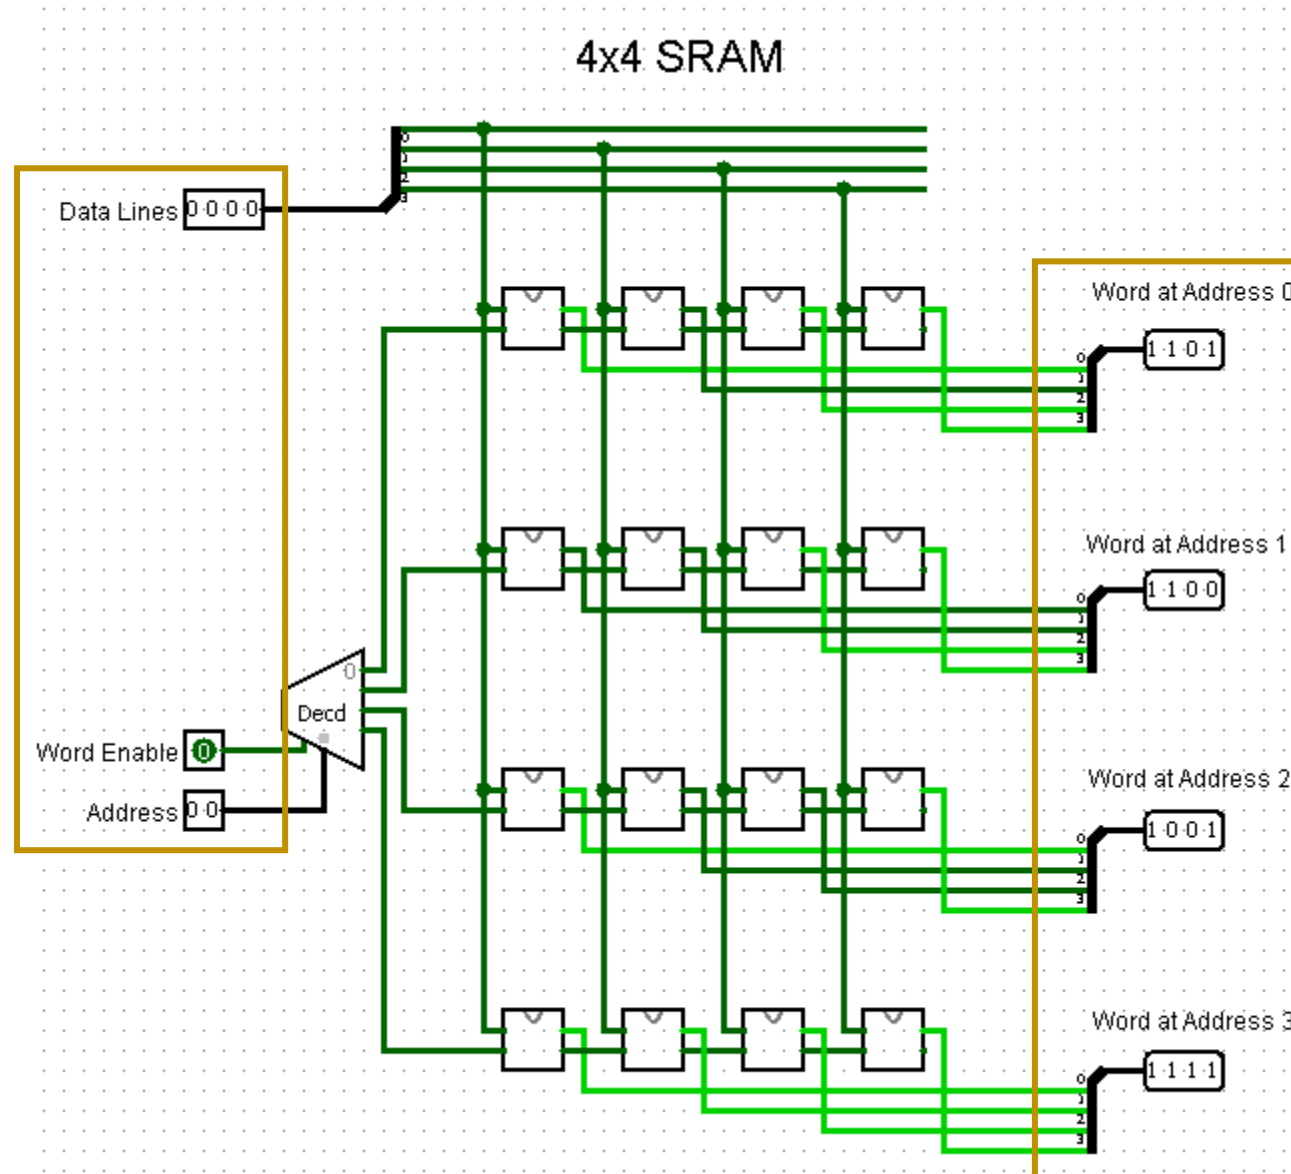
# RAM



# RAM



# RAM





# RAM

- Repeated reads and writes are required to refresh the capacitor in a DRAM cell
  - As opposed to the static storage of SRAM
  - This gives SRAM faster access than DRAM
- DRAM cells have fewer components and thus much smaller than SRAM cells
  - DRAM cells can be packed more densely than SRAM cells
  - DRAM is cheaper per bit than SRAM

# RAM

- An improvement of DRAM is **Synchronous DRAM (SDRAM)**
  - The clock keeps the processor and main memory synchronized.
- The use of a clock gives SDRAM the benefit of implementing pipelining
  - Pipelining allows a device to perform multiple operations at once
  - For example, SDRAM might output data to the processor while receiving the next address- simultaneously
  - We'll discuss pipelining in a later lecture

# RAM

- An improvement of SDRAM allows two words to be written or read during a single clock cycle
  - One word when the clock is 1, and another word when the clock is 0
- This is referred to as **Double Data Rate (DDR) SDRAM**
  - DDR2 doubles the data rate by allowing four words to be read/written during a single cycle
  - DDR3 quadruples the data rate by allowing eight words to be read/written during a single cycle
  - DDR4 (current technology)
    - Does not increase the words read/written during a single cycle
    - Uses a more advanced architecture for a higher transfer rate with a faster clock

# ROM

- **ROM (Read-Only Memory)** is a type of non-volatile memory that retains its stored bits.
- ROM is used by technologies that write data slowly, meaning writing data to the ROM is less frequent than reading data from it.
  - Contrary to what its name suggests, writing to ROM (*“programming the ROM”*) is possible.
- ROM commonly uses a floating-gate transistor (FGT), where electrons remain trapped even when no longer powered
  - A large positive voltage traps the electrons
  - A large negative voltage releases the electrons

# ROM

- There are several different types of ROM
  - **Masked-Programmed ROM:** The word line to bit line connections are hardwired and can never be changed
  - **One-Time Programmable ROM (OTP ROM):** The word line to bit line connections have a fuse that, when blown, can break the connection; Can only be programmed once.
  - **Erasable Programmable ROM (EPROM):** Electrons are trapped in FGTs using a large positive voltage; Electrons are released when the chip is exposed to ultraviolet light

# ROM

- **Electrically Erasable Programmable ROM (EEPROM):** Electrons are trapped in FGTs using a large positive voltage; Electrons are released from FGTs using a large negative voltage.
- **Flash:** A type of EEPROM; Electrons are trapped in FGTs using a large positive voltage; Using a large negative voltage, electrons in entire blocks of FGTs are released at once.



EPROM

# Cache Memory

- **Cache memory** is SRAM in the processor that holds:
  - The most frequently used data
  - The most recently accessed data
- Processors will often have several levels of cache memory.
  - **Level 1 (L1) Cache** – smallest, fastest cache
  - **Level 2 (L2) Cache** – larger, slower than the L1 cache
  - **Level 3 (L3) Cache** – larger, slower than the L2 cache
  - And so on...

# Cache Memory

- Cache memory holds copies of data from main memory in units called blocks
  - Each block has a fixed size of bytes
- For example, a cache memory could have 512 blocks that each store 128 bytes
  - $2^9 \text{ blocks} \times 2^7 \frac{\text{bytes}}{\text{block}} = 2^{16} \text{ bytes} = 64\text{KiB}$



# Cache Memory – Direct Mapping

- In a **direct mapped** cache, each block of main memory is mapped to a block in cache memory.
- Main memory is much larger than the cache, so multiple blocks of main memory will map to the same block in the cache

# Cache Memory – Direct Mapping

- To calculate the cache block that corresponds to a main memory block:

$$M_b \bmod B_c$$

- Where  $M_b$  is the block number of main memory block
- Where  $B_c$  is the total number of blocks in the cache

# Cache Memory – Direct Mapping

- When a word is to be *read/accessed* from main memory, the CPU first checks the cache to see if the block that contains the word is already in the cache.
  - If so, it simply obtains that data (**Cache Hit**)
  - If not, then the block is copied from main memory to its appropriate cache block (**Cache Miss**)

# Cache Memory – Direct Mapping


- When a word is to be *stored* to main memory, the CPU first checks the cache to see if that main memory address is already mapped in the cache.
  - If so, it simply stores that value in that block of the cache
  - If not, then the main memory block is copied to its appropriate cache block. The word is then written to that cache block.
- Either way, the data in the updated cache block must be copied back to main memory
  - Can be accomplished in one of two ways

# Cache Memory – Direct Mapping

- One way is that whenever a write operation occurs, the new data in the cache is written back to its corresponding address in main memory
- Another way is the cache stores a **dirty bit** (or **valid bit**) for each cache block
  - When the content of the block is changed, the valid bit is set to 1
  - When a new main memory block is to be copied to cache, the valid bit of the corresponding cache block is checked.
    - If 1, it writes the cache block to its corresponding main memory block, then loads the new main memory block (and resets the valid bit to 0)


# Cache Memory – Direct Mapping

- For illustration, here is a system with 32 bytes of main memory
  - Each square represents one byte
  - Each **column** is a block
  - Addresses 0 through 31



# Cache Memory – Direct Mapping

- We'll consider a simple CPU that has 16 bytes of cache memory
  - Each square represents one byte
  - Each **row** is a block
  - Addresses 0 through 7




# Cache Memory – Direct Mapping

- Let's see how the main memory address 10101 maps to the cache

- 10 1 01**


- 10 = **tag** or **main block**
- 1 = Cache block number
- 01 = Byte number in the cache block



MAIN MEMORY

tag	Cache block	Cache byte
10	1	01

RAM block




CACHE MEMORY



# Cache Memory – Direct Mapping


- The CPU is instructed to read the byte at address **10110** from main memory
  - The orange byte below



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
	1 00	
	1 01	
	1 10	
	1 11	

# Cache Memory – Direct Mapping


- First the CPU checks if the value is in the cache
  - **10 1 10**
  - Not present (*cache miss*)



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
0	1 00	
	1 01	
	<b>1 10</b>	
1	1 11	

# Cache Memory – Direct Mapping


- The block (block 5) is loaded into cache
  - 10 1 00
  - 10 1 01
  - 10 1 10
  - 10 1 11



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
10	1 00	10100
10	1 01	10101
10	1 10	10110
10	1 11	10111

# Cache Memory – Direct Mapping


- The CPU is instructed to read the byte at address 00101 from main memory
  - The yellow byte below



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
10	1 00	10100
10	1 01	10101
10	1 10	10110
10	1 11	10111

# Cache Memory – Direct Mapping

- First the CPU checks if the value is in the cache
  - 00 1 01**
  - Data in cache address 101, but wrong tag (*cache miss*)




Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
10	1 00	10100
10	1 01	10101
10	1 10	10110
10	1 11	10111

# Cache Memory – Direct Mapping

- The block (block 1) is loaded into cache


- 00 1 00
- 00 1 01
- 00 1 10
- 00 1 11



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111

# Cache Memory – Direct Mapping


- The CPU is instructed to read the byte at address **00100** from main memory



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	<div></div> 00100
00	1 01	<div></div> 00101
00	1 10	<div></div> 00110
00	1 11	<div></div> 00111

# Cache Memory – Direct Mapping


- First the CPU checks if the value is in the cache
  - **00 1 00**
  - Data in cache address 100 with correct tag (*cache hit*)
  - Does not retrieve from main memory








Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111



# Cache Memory – Direct Mapping


- The CPU is instructed to store a byte to address 00111 in main memory
  - Will be represented as 



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	 00100
00	1 01	 00101
00	1 10	 00110
00	1 11	 00111

# Cache Memory – Direct Mapping


- The CPU writes the value to the cache
  - 00 1 11
  - Data in cache address 111 with correct tag (*cache hit*)
    - Data in cache is updated



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111


# Cache Memory – Direct Mapping


- Data is stored to main memory (or dirty bit is set)



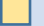



Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	<div></div> 00100
00	1 01	<div></div> 00101
00	1 10	<div></div> 00110
00	1 11	<div></div> 00111

# Cache Memory – Direct Mapping


- The CPU is instructed to store a byte to address 10001 in main memory
  - Will be represented as 







Tag	Cache Address Set Byte	Data/Main Memory Address
	0 00	
	0 01	
	0 10	
	0 11	
00	1 00	 00100
00	1 01	 00101
00	1 10	 00110
00	1 11	 00111

# Cache Memory – Direct Mapping

- First the CPU checks if the value is in the cache
  - **10 0 01**
  - No data in cache address 001 (cache miss)




Tag	Cache Address <i>Set Byte</i>	Data/Main Memory Address
	0 00	
	<b>0 01</b>	
	0 10	
	0 11	
00	1 00	 00100
00	1 01	 00101
00	1 10	 00110
00	1 11	 00111

# Cache Memory – Direct Mapping

- The block (block 4) is loaded into cache


- 10 0 00
- 10 0 01
- 10 0 10
- 10 0 11



Tag	Cache Address Set Byte	Data/Main Memory Address
10	0 00	10000
10	0 01	10001
10	0 10	10010
10	0 11	10011
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111

# Cache Memory – Direct Mapping


- The CPU writes the value to the cache
  - **10 0 01**
  - Data in cache is updated



Tag	Cache Address <i>Set Byte</i>	Data/Main Memory Address
10	0 00	10000
10	<b>0 01</b>	10001
10	0 10	10010
10	0 11	10011
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111

# Cache Memory – Direct Mapping

- Data is stored to main memory (or dirty bit is set)






















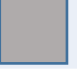
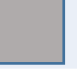


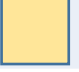
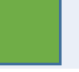
Tag	Cache Address <i>Set Byte</i>	Data/Main Memory Address
10	0 00	10000
10	<b>0 01</b>	10001
10	0 10	10010
10	0 11	10011
00	1 00	00100
00	1 01	00101
00	1 10	00110
00	1 11	00111



# Cache Memory – Direct Mapping

- Each row in the table below represents the read/write operations illustrated on the previous slides





Memory Address	Hit or Miss	Cache Contents							
		0	1	2	3	4	5	6	7
10110	M								
00101	M								
00100	H								
00111	H								
10001	M								

# Cache Memory – Associative Set Mapping

- In an **associative mapped** cache, each block of main memory is mapped to a *set of blocks* in cache memory.
- When a block is copied from RAM to the cache, the block can be stored in any one of the blocks that belong to the set.
- An ***n-way set associative cache*** indicates that each cache set contains  $n$  blocks per set.
  - A 3-way set associative cache has three blocks per set
  - A 1-way set associative cache has one block per set
    - A 1-way set associative cache *is* direct mapping

# Cache Memory – Associative Set Mapping


- 2-way associative cache
- Consider the gray square/byte with the address **10010**



# Cache Memory – Associative Set Mapping



- **10 0 10**

- 10 = **tag** or **main block**
- **0** = Cache set number
- 10 = Byte number in the cache block




# Cache Memory – Associative Set Mapping

- The CPU is instructed to read the byte at address **10101** from main memory
  - The orange byte below





# Cache Memory – Associative Set Mapping

- **10 1 01**



# Cache Memory – Associative Set Mapping

- Within this set, there are two blocks to choose from (blocks 2 and 3).
  - Which block is chosen will depend on the *block replacement strategy* implemented




# Cache Memory – Associative Set Mapping

- One such strategy is called **Least Recently Used (LRU)**
  - The block to be replaced in the set is the block that has been sitting, unused, for the longest time.
- Another block replacement strategy is called **First In, First Out (FIFO)**
  - The block to be replaced in the set is the block that has been in the cache for the longest time.
- A third strategy is to randomly choose a block for replacement
- Each strategy has their pros and cons
  - No perfect/optimal block replacement strategy



# Cache Memory – Associative Set Mapping

- We'll assume LRU.
  - Block 2 is chosen





# Cache Memory – Associative Set Mapping

- **00101**

tag	Cache set	Cache byte
00	1	01
RAM block		



Tags don't match (miss)



# Cache Memory – Associative Set Mapping

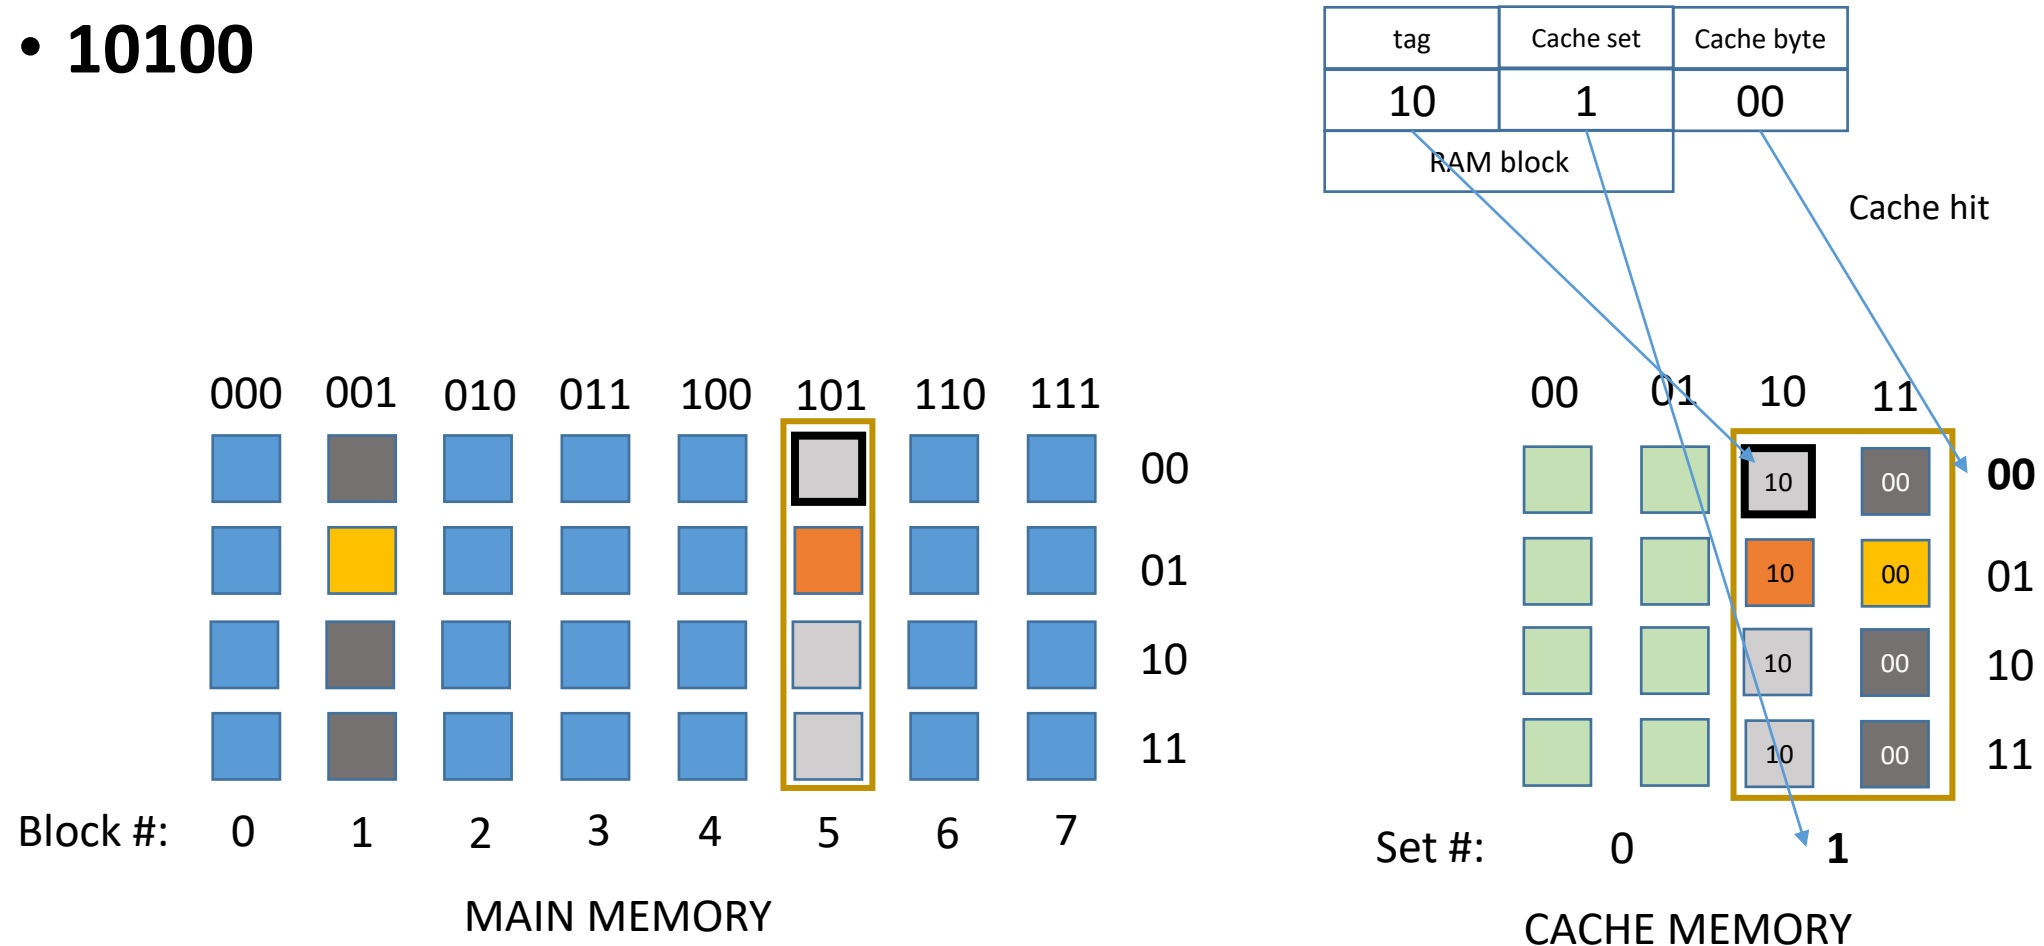
- **00101**

tag	Cache set	Cache byte
00	1	01
RAM block		




# Cache Memory – Associative Set Mapping

- **10100**




# Cache Memory – Associative Set Mapping

- **01100**




# Cache Memory – Associative Set Mapping

- Since we are using LRU, we will replace block 3 (since we just used block 2)
- Had we been using FIFO, then block 2 would have been replaced because that has been in the cache longer than the data in block 3



# Cache Memory – Associative Set Mapping



- Each row in the table below represents the read/write operations illustrated on the previous (associative mapping) slides



Memory Address	Hit or Miss	Cache Contents							
		Set 0 (Block 0)	Set 0 (Block 1)	Set 1 (Block 2)		Set 1 (Block 3)			
10101	M			10	10	10	10		
00101	M			10	10	10	10	00	00
10100	H			10	10	10	10	00	00
01100	M			10	10	10	10	01	01

# Cache Memory – Fully Associative Set Mapping

- A **fully associative mapped** cache consists of one set that contains every block of cache memory.







# Cache Memory – Fully Associative Set Mapping

- **00000**



tag	Cache set	Cache byte
00	0	00
RAM block		



# Cache Memory – Fully Associative Set Mapping

• **00100**



tag	Cache set	Cache byte
00	1	00
RAM block		



# Cache Memory – Fully Associative Set Mapping

- **01000**



tag	Cache set	Cache byte
01	0	00
RAM block		



# Cache Memory – Fully Associative Set Mapping

- **01100**

tag	Cache set	Cache byte
01	1	00
RAM block		





# Cache Memory – Fully Associative Set Mapping

- **10000**


- Out of space
- The block to replace (LRU or FIFO in this case) is block 0

tag	Cache set	Cache byte
10	0	00
RAM block		



# Cache Memory – Associative Set Mapping

- Each row in the table below represents the read/write operations illustrated on the previous (associative mapping) slides



Memory Address	Hit or Miss	Cache Contents															
		Block 0				Block 1				Block 2				Block 3			
00000	M	00	00	00	00												
00100	M	00	00	00	00	00	00	00	00								
01000	M	00	00	00	00	00	00	00	00	01	01	01	01				
01100	M	00	00	00	00	00	00	00	00	01	01	01	01	01	01	01	
10000	M	10	10	10	10	00	00	00	00	01	01	01	01	01	01	01	